

DAZLARUS INDUSTRIES

PlanSpec

An open standard for structured AI execution plans.

What

An **open standard (MIT)** for structured AI execution plans. When an AI agent needs to do real work — write code, modify infrastructure, operate systems — PlanSpec defines the format for the plan that gets reviewed, approved, and executed.

Why It Matters

AI agents are moving from chat to execution. Today, the "plan" is an unstructured LLM response — no accountability, no validation, no audit trail. PlanSpec fixes that.

How It Works

A PlanSpec document captures:

- **What** the AI will do (scoped steps with dependencies)
- **Why** this approach (recorded rationale and alternatives considered)
- **What's off-limits** (explicit scope exclusions)
- **What could go wrong** (risks with mitigations)
- **What requires human approval** (delegation boundaries)
- **How we know it's done** (measurable success criteria)

Every plan is versioned, revision-tracked, and machine-validatable.

Version	1.2.0
License	MIT
Format	JSON (canonical), Markdown (rendered view)
Schema	JSON Schema Draft 07
Repository	github.com/Dazlarus-Industries/planspec

Who It's For

- **Engineering teams** deploying AI agents for code, infrastructure, or operations
- **Security/compliance** teams who need oversight and audit trails for AI-driven work
- **Tool builders** who want a proven plan format instead of inventing their own
- **Enterprise architects** defining governance policies for AI agent deployment

Get Involved

- Read the spec: [SPEC.md](#)
- Try the examples: [examples/](#)
- Open an issue with feedback or questions
- Reach out: darien@dazlarus.dev